

BOOTS ON THE GROUND AI

AI Solutions for Small Business Owners

CLAUDE PLAYBOOK

Pick the Right Model for Every Job

Your owner-friendly guide to choosing between Claude's three models: Haiku, Sonnet, and Opus. Task-by-task recommendations, plain-English tradeoffs, and the prompts to try today.

Built for small-business owners who want AI that pulls its weight — without burning through rate limits or paying for horsepower they don't need.



CLAUDE



BOOTS APPROVED



TABLE OF CONTENTS

Your roadmap to getting the most out of Claude

01	Why Model Choice Matters (The Boots Philosophy)	3
02	Meet the 3 Models	4
03	The 7 Rules of Model Selection	5
04	Task-by-Task Recommendations	7
05	Pricing at a Glance	9
06	Claude.ai vs Claude Code vs API	10
07	Cost-Saving Tricks (Caching & Batch)	11
08	Common Mistakes Small Teams Make	12
09	Model Selection Checklist (Tear-Out)	13



01 WHY MODEL CHOICE MATTERS

The Boots Philosophy on Picking AI

If you've ever asked an AI to do a quick lookup and watched it chew through your rate limit for five minutes to give you a three-word answer — or asked a cheap fast model to tackle a deep analysis and gotten back confident nonsense — the problem isn't Claude. It's using the wrong tool for the job.

BOOTS SAYS:

Think of Claude's three models like a three-tool toolbox. Haiku is a fast multi-bit screwdriver — perfect for 90% of quick jobs. Sonnet is the cordless drill — your daily workhorse. Opus is the rotary hammer — absolutely the right tool for concrete, absolute overkill for hanging a picture frame.

Without good model selection, you face:

- ✗ Burning through rate limits on tasks a cheaper model would handle
- ✗ Paying 5x more per task than necessary (Opus costs 5x Haiku)
- ✗ Waiting 30 seconds for an answer that should take 2
- ✗ Underpowered answers when you actually needed deep reasoning
- ✗ Guessing every single time instead of having a default

With good model selection, you get:

- + Fast answers on fast tasks, deep answers on hard ones
- + Rate limits that actually cover your workday
- + Predictable costs you can budget against
- + Better output quality because the right model sees the task
- + A playbook your whole team can follow without thinking

THE BOOTS STANDARD:

Every Boots On The Ground AI workflow specifies which Claude model to use for each task. No 'just try whichever'. If you can't justify the choice, you're probably overpaying or underperforming. Pick on purpose.



02 MEET THE 3 MODELS

Haiku, Sonnet, Opus — in plain English

Anthropic ships three current-generation Claude models. They all speak the same way, handle images, and follow instructions — but they differ dramatically in speed, cost, and depth of reasoning.

Side-by-Side Comparison

Model	Speed	Best For	API Cost (in/out per 1M tokens)
Haiku 4.5	Fastest	Quick answers, summaries, extraction, routing, high-volume	\$1 / \$5
Sonnet 4.6	Fast	Writing, coding, analysis, everyday business work	\$3 / \$15
Opus 4.7	Slower	Deep research, complex coding, multi-step reasoning	\$5 / \$25

BOOTS EXPLAINS:

Notice the 5x spread. Haiku at \$1/M input is five times cheaper than Opus at \$5/M. Output pricing keeps the same ratio. If you run Opus on a task Haiku could handle, you pay five times more for zero gain — and you burn five times more of your rate limit.

The 30-Second Description

1 Haiku 4.5 — The Sprinter

Lightning-fast, cheapest, handles simple questions and short answers instantly. Rate-limit-friendly. Use it whenever you don't specifically need more.

2 Sonnet 4.6 — The Workhorse

Anthropic's recommended default and the model most small-business owners should reach for daily. Strong at writing, coding, analysis, and multi-step workflows. Handles 80% of real work beautifully.



3

Opus 4.7 — The Specialist

The flagship. Use it when Sonnet has struggled, or when the stakes are high enough that quality matters more than cost. Long research reports, hard debugging, dense technical analysis. Reserve for tasks that need it.



03 THE 7 RULES OF MODEL SELECTION

Memorize these. They replace guessing with a decision.

1 Start With Sonnet

When in doubt, default to Sonnet. It's fast enough for interactive work, smart enough for 80% of real business tasks, and priced in the middle. Most workflows never need anything else.

2 Drop to Haiku for Simple, Repetitive Tasks

Extracting phone numbers from messages. Summarizing a single email. Classifying support tickets into buckets. If the task is mechanical and the answer is short, Haiku does it for 1/3 the cost at 3x the speed.

3 Climb to Opus for Hard Reasoning

Diagnosing a bug Sonnet missed. Drafting a contract. Comparing 20 vendors across 6 criteria. If you've tried Sonnet and the output fell short, or the cost of a mistake is high, that's your Opus moment.

4 Match the Model to the Output Length

Long, detailed outputs (1,000+ words) on complex topics benefit from Opus's deeper reasoning. Short, focused outputs rarely do. Length alone isn't the trigger — complexity is.

5 Let Rate Limits Guide Routine Choices

If you're hitting rate limits on Claude.ai, you're probably using a heavier model than needed. Shift routine tasks to Haiku or Sonnet so Opus is available when it counts.

6 Test Before You Commit

A new Claude version isn't a patch on the old one — each release is a separate training run. A task that suited Opus 4.5 might sit more naturally on Sonnet 4.6. When a new model ships, spend 10 minutes re-testing your go-to prompts.



7

Document Your Defaults

Write down which model you use for which recurring task. 'Sonnet for blog drafts, Haiku for customer-reply classification, Opus for quarterly reports.' Stop re-deciding every time.



04 TASK-BY-TASK RECOMMENDATIONS

The task, the pick, and why

Below is a task-by-task cheat sheet for the work most small-business owners actually do. Pick your task, pick your model, get to work.

Writing & Content

Task	Pick	Why
Draft a blog post	Sonnet	Balances quality and cost; handles tone and structure well.
Polish an already-written draft	Haiku	Editing is mostly mechanical; Haiku is fast and cheap.
Ghostwrite a thought-leadership piece	Opus	Needs voice, nuance, and original insight Sonnet can miss.
Summarize an article or email	Haiku	Straightforward extraction — no reasoning needed.
Write 10 social media post variants	Sonnet	Needs creativity and tone control at modest scale.
Translate a document	Sonnet	Reliable across major languages; cheaper than Opus.

Customer Service & Email

Task	Pick	Why
Classify support tickets (bug vs billing vs feature)	Haiku	Classic classification — volume is high, complexity is low.
Draft a standard customer reply	Sonnet	Needs context and warmth; Haiku can feel flat.
Respond to a complaint that could escalate	Opus	Stakes are high; nuance matters; reputation is on the line.
Extract order numbers from email threads	Haiku	Pure extraction — minimum viable reasoning.
Auto-generate FAQ answers from docs	Sonnet	Needs to understand context and rephrase accurately.

Coding & Technical

Task	Pick	Why
------	------	-----



Write new code (typical feature)	Sonnet	Anthropic's default for Claude Code — strong on most everyday work.
Debug a stubborn bug	Opus	Deep reasoning finds issues Sonnet skips — race conditions, logic traps.
Scaffold boilerplate / UI components	Haiku	Fast, repetitive pattern work where Opus is overkill.
Review and refactor a module	Opus	Catches missed dispose calls, async bugs, dead code paths.
Generate SQL from plain English	Sonnet	Needs real schema understanding but rarely deep reasoning.
Comment undocumented code	Haiku	Mechanical pass over already-written code.

Research & Analysis

Task	Pick	Why
Quick factual lookup	Haiku	Small question, short answer — maximum efficiency.
Compare 3 products or vendors	Sonnet	Structured comparison that doesn't need deep synthesis.
Synthesize 10+ sources into a report	Opus	Sustained reasoning across long inputs — Opus's sweet spot.
Analyze a spreadsheet or data file	Sonnet	Handles tabular reasoning well without Opus pricing.
Forecast or scenario-plan from messy data	Opus	Multi-step reasoning where small errors compound.

Business Operations

Task	Pick	Why
Meeting notes → action items	Sonnet	Needs to distinguish decisions from discussion.
Convert handwritten receipt photos to text	Haiku	Vision + extraction, no reasoning required.
Draft a contract or SOW	Opus	Legal language, stakes, and nuance all favor the flagship.
Build a slide deck outline	Sonnet	Creative structuring at middle-tier cost.



Tag hundreds of inventory items

Haiku

Classification at volume — Haiku economics are mandatory.

BOOTS TIP:

A quick self-test: if you could describe the task in under 10 seconds and the output is under 100 words, Haiku probably handles it. If the task requires you to think hard before you even write the prompt, reach for Opus.



05 PRICING AT A GLANCE

What you actually pay, and where

Two ways to use Claude: via the Claude.ai subscription (flat monthly), or via the API (pay-per-use by tokens). Most small-business owners use the subscription; teams integrating Claude into software use the API.

Claude.ai Subscription Plans

Plan	Price	Models Included	Best For
Free	\$0	Haiku, Sonnet	Trying Claude, light personal use
Pro	\$20/month	Haiku, Sonnet, Opus	Solo owners using Claude daily
Max	\$100–\$200/month	All models, higher limits	Power users hitting Pro rate limits
Team	\$25+/user/month	All models, admin & sharing	Small teams (2+ people)

API Pricing (per Million Tokens)

Model	Input	Output	Context Window
Haiku 4.5	\$1	\$5	200K tokens
Sonnet 4.6	\$3	\$15	1M tokens
Opus 4.7	\$5	\$25	1M tokens

BOOTS EXPLAINS:

A token is roughly 3/4 of a word in English. A 1,500-word article is about 2,000 tokens. A typical back-and-forth customer service chat might use 500 input + 200 output tokens. At Sonnet rates, that's about \$0.005 per chat — \$5 per thousand chats.



06 CLAUDE.AI VS CLAUDE CODE VS API

Three ways to use Claude — pick the right one

Anthropic offers three different products, each suited to different kinds of work. Most small-business owners only need the first; builders reach for the other two.

1 Claude.ai (the chat app)

The web and mobile chat interface at claude.ai. Great for daily writing, research, analysis, and general-purpose AI work. Flat monthly fee. Use this for 90% of non-coding small-business work.

2 Claude Code (the terminal tool)

A command-line tool for developers. Lets Claude directly edit files, run commands, and build software on your machine. Included free in Pro and above. Use this if you're writing code yourself.

3 Claude API (build it into your product)

Pay-per-token access for integrating Claude into your own software. Use this if you're shipping AI features in your app, running automated pipelines, or processing documents at scale.

BOOTS SAYS:

If you just want to use Claude as a smart assistant for your business, stop at Claude.ai Pro (\$20/mo). Everything else exists for developers building with Claude. Don't let anyone tell you you need the API for normal work.



07 COST-SAVING TRICKS

How to cut your API bill by up to 95%

If you're using the API at volume, two built-in discounts can dramatically reduce your costs. Both stack with each other.

Trick 1: Prompt Caching (up to 90% off)

If you send the same system prompt, document, or context with every request, Anthropic caches it for you after the first call. Cached input costs 10% of the normal rate. On a Sonnet workflow that repeats a 2,000-token system prompt across 1,000 daily requests, caching can cut your monthly input cost by 50% or more.

prompt_caching.py

```
# Mark the cached content with cache_control
messages = [{
    "role": "system",
    "content": [{
        "type": "text",
        "text": long_system_prompt,
        "cache_control": {"type": "ephemeral"}
    }]
}]
```

Trick 2: Batch API (50% off across the board)

If you don't need answers immediately — overnight data processing, nightly report generation, backfilling classifications — submit requests via the Batch API. You get answers within 24 hours for exactly half the token cost. Haiku drops from \$1/\$5 to \$0.50/\$2.50. Sonnet drops from \$3/\$15 to \$1.50/\$7.50. Opus drops from \$5/\$25 to \$2.50/\$12.50.

BOOTS EXPLAINS:

Caching + batch combined can bring effective costs down by 95%. If you're processing 1M customer emails a month on Haiku with a cached system prompt via batch, you're paying roughly \$25 for the whole month — not \$200.

SECURITY ALERT:

Never paste passwords, API keys, credit cards, or regulated data (HIPAA-PHI, financial PII) into the free Claude.ai or consumer tier. Use the Enterprise or Team plans with a signed BAA (for HIPAA) and the API with proper data handling for anything regulated. When in doubt, treat it as sensitive.



08 COMMON MISTAKES

What small teams get wrong

1 Using Opus by Default

The most common mistake. Opus feels 'smarter' so people default to it — and pay 5x more on every task. Sonnet handles 80% of work. Start there.

2 Using Haiku for Nuanced Writing

The opposite mistake. Haiku is fast and cheap but it can feel flat on tone-sensitive copy. If your email sounds like a robot, try Sonnet.

3 Never Re-Testing When New Models Ship

A task that suited Sonnet 4.5 might be better on Haiku 4.5 today. When Anthropic releases a new version, spend 10 minutes re-running your top 3 prompts and comparing output quality.

4 Pasting Sensitive Data Into Free Accounts

The free tier is fine for personal work. But PHI, client financials, and trade secrets need the Team/Enterprise plans with proper data controls and BAAs.

5 Ignoring Caching on Repeated Context

If you're running the same system prompt across thousands of API calls and not using prompt caching, you're leaving 50-90% of your input cost on the table.

6 Trusting Made-Up Statistics

All models occasionally hallucinate specific numbers, quotes, and links — even Opus. Always verify any stat, study, or URL before publishing AI-generated content.



09 MODEL SELECTION CHECKLIST

Print this. Pin it next to your monitor.

BOOTS APPROVED

BOOTS ON THE GROUND AI CLAUDE MODEL SELECTION CHECKLIST. Run through this before starting any new Claude workflow. 60 seconds of thought saves hours of rework and dollars of wasted tokens.

QUICK DECISION

- ☐ I can describe this task in under 10 seconds → try Haiku first
- ☐ This is typical daily business work → use Sonnet (the default)
- ☐ Sonnet already struggled, or stakes are high → escalate to Opus
- ☐ This is high-volume and mechanical → Haiku + Batch API

BEFORE I RUN THE PROMPT

- ☐ I picked a model on purpose, not by habit
- ☐ I've checked the output is short enough that this model makes sense
- ☐ If API: I enabled prompt caching if the system prompt repeats
- ☐ If API and not urgent: I considered Batch for 50% off
- ☐ No passwords, client data, or regulated info in the prompt

AFTER THE FIRST OUTPUT

- ☐ Output meets my quality bar (if not, try a stronger model)
- ☐ Output didn't take forever (if so, try a lighter model)
- ☐ I didn't burn unnecessary rate limit
- ☐ I verified any stats, quotes, or links before publishing
- ☐ I noted this task's working model in my personal playbook

QUARTERLY REVIEW

- ☐ Re-tested top 3 prompts against the latest model releases
- ☐ Checked if any tasks have drifted to a different model's sweet spot
- ☐ Reviewed monthly API costs for obvious overspend
- ☐ Trained team on current model defaults